

Amazon Product Co-purchasing Analysis

Hao Wang, Tianxiang Wang, Tianhao Zhou

Mar 17, 2025

Introduction

- The rise of e-commerce platforms like Amazon has created vast product networks with complex co-purchasing relationships. Traditional databases struggle with these interactions.
- Our project integrates the Amazon Meta Dataset into a hybrid SQL + Graph + NoSQL system to analyze product influence, enhance recommendations, and leverage graph-based approaches for deeper insights.



Amazon Product Recommendation System

♦ What is the Amazon product recommendation system?

- The Amazon recommendation system suggests products based on customer behavior, past purchases, and purchasing history.

♦ Why do we need the Amazon recommendation system?

- It helps customers find products easily, improves their shopping experience, and increases sales.

♦ How does the Amazon recommendation system work?

- Product data is stored in SQL, co-purchasing patterns are analyzed in graphs, and fast recommendations are cached in NoSQL.

Dataset-amazon products

Amazon Product Co-Purchasing Network Metadata

◆ Source:

- Data crawled from the Amazon website during the summer of 2006.

◆ Structure & Content:

- **Product Information:** Title, ASIN, and product group (e.g., Books, Music CDs, DVDs, Videos).
- **Sales Data:** Amazon sales rank.
- **Network Connections:** List of similar products (i.e., products that are co-purchased).
- **Categorization:** Detailed hierarchical product categories with corresponding category IDs.
- **Review Information:** Includes review date, customer identifier, rating, number of votes, and helpfulness count.

◆ Role in Application:

- Establishes a co-purchasing network among products.
- Supports market analysis and evaluation of product performance.

Relational Database - Pre Processing

Explored relational dataset(PostgreSQL): Identified columns (**ASIN**, **Title**, **Product Group**, **Sales Rank**, **Categories**, **Similar Products**).

♦ Data Cleaning:

- Removed **duplicates**.
- Handled **missing values** (e.g., filtering out NULL titles, categories).
- **Formatted categories** for easier querying.
- **Excluded invalid rankings (-1)** in queries(Discontinued Products).

♦ Schema Design in PostgreSQL:

- Created **amazon_products** table for product metadata.
- **Indexed ASIN** for faster lookups.

♦ Imported data using COPY command:

- Loaded CSVs into PostgreSQL (**\COPY** used for permission issues).

Relational Database

◆ Key Queries to gain insights of the data:

🔍 Find top-selling book:

```
--find best selling products--
SELECT title, sales_rank
FROM amazon_products
WHERE sales_rank > 0 -- Excludes -1 (invalid rankings)
ORDER BY sales_rank ASC
LIMIT 10;
```

title	sales_rank
The War of the Worlds	1
Shirley Valentine	2
Leslie Sansone - Walk Away the Pounds - Super Fat Burning	6
Robin Hood - Men in Tights	7
Richard Simmons - Sweatin' to the Oldies	8
Howard the Duck	12

🔍 Find products with many similar co-purchases:

```
--Find products with high purchasing rate
SELECT asin, title,
       LENGTH(similar_products) - LENGTH(REPLACE(similar_products, '|', '')) + 1 AS similar_count
FROM amazon_products
WHERE similar_products IS NOT NULL AND similar_products <> ''
ORDER BY similar_count DESC
LIMIT 10;
```

	asin	title	similar_count
1	1893115542	Data Mining & Statistical Analysis Using SQL	5
2	0691020841	The Camphor Flame	5
3	091912366X	The Mysterium Lectures: A Journey Through C.G. Ju...	5
4	0312981252	Bad Boy: The Murderous Life of Kenneth Allen McDu...	5
5	0064420280	Stanley and the Magic Lamp	5

🔍 Find products with low co-purchase rate:

```
--Find Products with Low Co-Purchases--
SELECT asin, COUNT(*) AS co_purchase_count
FROM amazon_products
WHERE similar_products IS NOT NULL
GROUP BY asin
ORDER BY co_purchase_count ASC;
```

	asin	co_purchase_count
1	1569715076	1
2	0963290614	1
3	0802773613	1
4	0806522178	1
5	0715307975	1

Graph Database- Data Preprocessing

Relationships

Product – CO_PURCHASED_WITH – Another Product

Product – BELONGS_TO – Category

```
SELECT asin AS product_asin,
       UNNEST(string_to_array(similar_products, '|')) AS similar_asin
FROM   amazon_products
WHERE  similar_products IS NOT NULL
       AND similar_products <> '';
```

```
( asin,      similar_asin )
( <that product>, 0804215715 )
( <that product>, 156101074X )
( <that product>, 0687023955 )
```

similar_products
1559360968 1559361247 1559360828 155936101...
1585741485 0140246967 1557504288 037420518...
0071410546 1580531784 1578200326 B0000A2W55
B000059QC1 B00000JQIE B00029J1X6 B0006TR06...
0785114572 078511078X 0785114033 078511404...
0865778973 0071343547 0721662854 072168173...
B00007JMD8 6305350221 B00004RF9B B00005JKF...
B00000JCDS B000004CSZ B00016XN6Q B00005LLY...

- Multiple IDs for the same category name
- Overly generalized categories

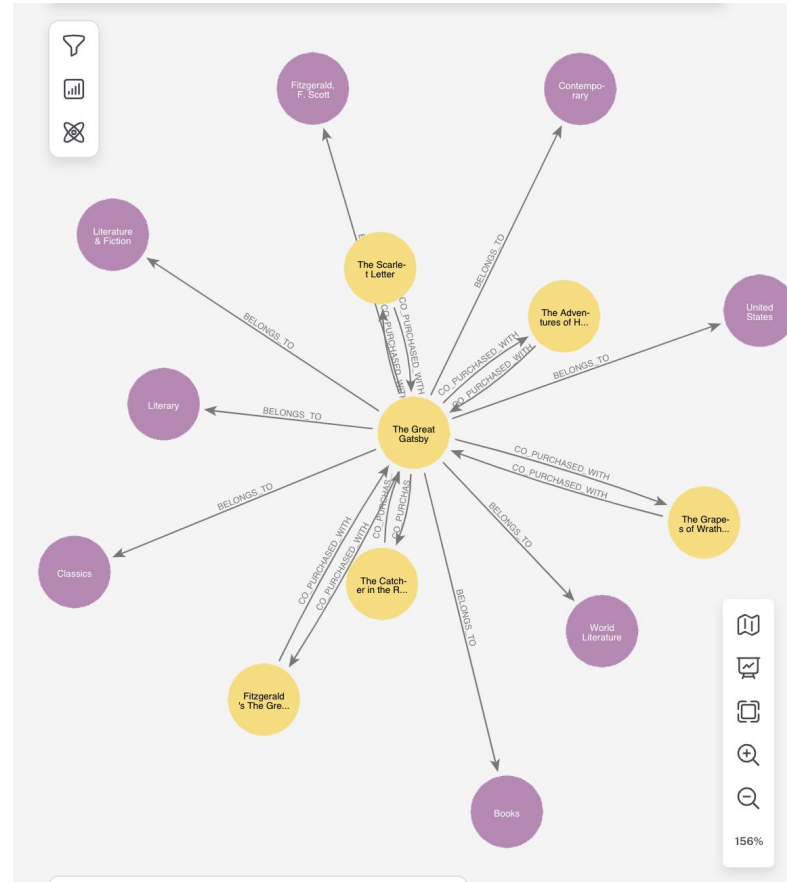
-Redundancy

–Missing a clear delimiter

Graph Database

PageRank

Louvain Community Detection



Key-value Database

♦ Objective:

- Store frequently accessed product attributes in Redis for rapid lookups, reducing MongoDB load.

♦ Implementation:

- **Cached Attributes:** Product metadata (ASIN, title, category, sales rank, reviews).
- **Key Patterns:**
 - `top_products:{category}:limit_{limit}`: Top-selling products (e.g., 10 best by sales rank).
 - `product:{asin}`: Individual product details (e.g., "0738700797").
 - `most_reviewed:limit_{limit}`: Products with most reviews (e.g., Harry Potter editions).
 - `promoted:{asin}`: Promoted products (e.g., "Candlemas: Feast of Flames").

♦ Results:

- **Sample Data:**
 - Top Sales: "Vehicular Technology..." (ASIN: 0780357213).
 - Most Reviewed: "Harry Potter..." (ASIN: 043936213X, 4995 reviews).
 - Promoted: "Candlemas..." (ASIN: 0738700797).

```
### 📌 Promoted Products Handling ###  
def add_promoted_product(asin, promotion_reason):  
    """Add a product to the promoted list"""  
  
    try:  
        product = products_collection.find_one({"asin": asin})  
        if not product:  
            print(f"❌ Product {asin} not found!")  
            return  
        promoted_data = {  
            "asin": asin,  
            "title": product["title"],  
            "reason": promotion_reason,  
            "added_date": "2025-03-18" # Use datetime.now().isoformat() in production  
        }  
        promoted_collection.update_one({"asin": asin}, {"$set": promoted_data}, upsert=True)  
        r.set(f"promoted:{asin}", json.dumps(promoted_data), ex=7*24*3600)  
        print(f"✅ Promoted product added: {asin}")  
    except Exception as e:  
        print(f"❌ Error adding promoted product: {e}")
```

Combine Databases

♦ Goal:

- Use **PostgreSQL** (structured metadata) + **Neo4j** (graph-based recommendations) + **MongoDB** (fast lookup for promotions)
- Build an **intelligent hybrid recommendation system**.

♦ Connection Setup

Python integration with:

- **PostgreSQL** (**psycopg2**) for structured data.
- **Neo4j** (**neo4j-driver**) for graph-based recommendations.
- **MongoDB** (**pymongo**) for quick promotional product retrieval.
- **Figure on the right** shows the connection setup.

♦ How It Works:

Step 1: Query **Neo4j** to find **co-purchased items** (graph relationships).

Step 2: Query **PostgreSQL** to get **product details and user ratings**.

Step 3: Query **MongoDB** to retrieve **promotional products for fast lookup**.

Step 4: Combine results, rank by **sales rank & average rating**, and return recommendations.

Database Connection and Check

```
# PostgreSQL (change to your own database/user/password/host/port)
pg_conn = psycopg2.connect(
    database="amazon_db",
    user="postgres",
    password="l01123",
    host="localhost",
    port="5432"
)
pg_cursor = pg_conn.cursor()

# Test PostgreSQL
pg_cursor.execute("SELECT COUNT(*) FROM amazon_products;")
result = pg_cursor.fetchone()
print("PostgreSQL is connected, total products:", result[0])

# Neo4j
neo4j_driver = GraphDatabase.driver("bolt://localhost:7687", auth=("neo4j", "2000l011"))

# Test Neo4j
def test_neo4j_connection():
    query = "MATCH (p:Product) RETURN count(p) AS total_products"
    with neo4j_driver.session() as session:
        result = session.run(query)
        total_products = result.single()[0]
        print("Neo4j is connected, total products:", total_products)

test_neo4j_connection()

# MongoDB
mongo_client = pymongo.MongoClient("mongodb://localhost:27017/")
db = mongo_client["amazon"]
promoted_collection = db["promoted_products"]

# Redis
r = redis.Redis(host='localhost', port=6379, db=0)

PostgreSQL is connected, total products: 548552
```

DEMO

Lessons Learned

◆ Version 1: Exact Matching Issues

- Only returned products with **exact title match**, missing relevant results.
- Solution: Switched to **partial title search**.

◆ Version 2: Basic Co-Purchase Recommendations

- Recommended **products directly linked** in Neo4j.
- Problem: Lacked **flexibility and ranking** for better results.

◆ Version 3: Smarter Recommendations with Multiple Factors

- Introduced **ranking by rating** (from PostgreSQL).
- Improved Neo4j recommendations with **community-based product associations**.

◆ Version 4: Fully Integrated Hybrid System

- Combined **SQL, Graph, and NoSQL** for a **multi-source recommendation**.
- Ensured **fast lookup, category-based, and community-based** recommendations.
- Balanced **speed, relevance, and scalability**.

◆ Key Takeaways

- **Data preprocessing** is critical for large datasets.
- **Hybrid systems** provide more accurate and diverse recommendations.
- **Database integration** requires careful optimization to avoid performance bottlenecks.
- **Real-world recommendation systems** need flexibility in **search queries and ranking factors**.

Conclusion

Contributions

- Successfully integrated **PostgreSQL, Neo4j, and MongoDB** into a unified recommendation system.
- Built a **hybrid recommendation model** combining relational, graph, and key-value store techniques.
- Optimized data processing by **preprocessing categories** and reducing redundant data.

Challenges Overcome:

- **Large-scale data issues** (memory constraints in Neo4j).
- **Balancing accuracy vs. efficiency** in recommendation ranking.

Future Directions:

- **Improve recommendation logic** (e.g., collaborative filtering, deep learning-based ranking).
- **Enhance user personalization** (e.g., user behavior-based recommendations).
- **Optimize query performance** (e.g., indexing, caching for faster lookups).
- **Detects fake reviews and suspicious activities**, ensuring trust and better recommendations on the platform.

Thank you!